

The Beauty of Linux I

by Imad Rahal

A. **Redirecting Output:** Every Linux process has three file descriptors:

- *standard in **stdin*** used as the source of input for the process (e.g. keyboard) ... **file descriptor 0**
- *standard out **stdout*** used as the destination for process's output (e.g. screen) ... **file descriptor 1**
- *standard error **stderr*** used as the destination for process's errors messages (e.g. screen) ... **file descriptor 2**

Processes are written generically to read from *keyboard* as **stdin** and write to *screen* as **stdout** but that can be changed. Redirecting I/O can be accomplished on Linux via **<** and **>**, respectively (**>&** for **stderr**). E.g.:

- (1) Command **ls** lists the contents of the current folder. Open a Linux terminal and try it out.
- (2) Now, issue command **ls > my.file** [this redirects the output of the **ls** command to a new file called **my.file**] and then issue command **cat my.file** to see the contents of the file [cat displays the contents of the specified file in the terminal]
- (3) Issue command **sort** and hit enter followed by a number of text lines to sort (when done entering data, press CTRL-d) to see how it works. For example, try the following lines:
once upon a time
a small creature
came to live in
the forest
^D
a small creature
came to live in
once upon a time
the forest
[This **sort** command sorted the lines!]
- (4) Let us repeat the above by issuing command **sort > story.txt** first followed by the text data from above (press CTRL-d when done) and then issue command **cat story.txt**
- (5) Now, place the same data to sort in a file called **story2.txt** and issue command **sort < story2.txt**

B. **Linux Pipes:** Output of one Linux command can be sent as input to another making a connection known as a *pipe*. Linux pipes are created using the **|** character (PS: Linux command **wc** prints the number of *lines*, *words*, and *bytes* in a given file)

- (1) issue command **ls | wc > wc.fields** [this example forwards the output of the **ls** Linux command to the **wc** command which counts the number of lines, words and bytes and sends output to file **wc.fields**].
- (2) Use the **cat** command to see the contents of file **wc.fields**

C. **File Types & Contents:** In Linux, file names are made up of alphanumeric characters, periods and underscores (BTW, a dot or period in a file name means nothing in Linux! Extensions aren't recognized on Linux). Names are case sensitive. To view the contents of a given file, the following commands can be used:

- **cat filename**
- **more filename** (used with large files – use space bar to go thru file)
- **head-x filename** (displays the first **x** lines in the file)
- **tail-x filename** (displays the last **x** lines in the file)
- (1) Try the above commands using file **story2.txt** created earlier in A.

The command **file filename** looks at the first several hundred bytes of a given file (i.e. **filename**) and does a statistical analysis of the types of characters that it finds there.

- (2) issue command **file wc.fields**
- (3) issue command **file** on a **.java** and a **.class** file

D. Permissions: Permissions in Linux are divided into three categories: *User* (or Owner), *Group* and *Others*. Any file belongs to a single user and thus to a single group. Files have separate *Read/Write/Execute* permissions for each of the above categories in this order: *User* permissions, *Group* permissions and then *Others* permissions. E.g. **drwxr-xr--** means that this is a directory not a file (starts with a *d*) where the owner can **r** (i.e. read), **w** (write) and **x** (execute), group members can **r** and **x** while others can only **r**).

- (1) Use **ls -l** to view current permissions on files and folders in your home directory.
- Command **chmod who=permissions filename** changes the permissions on a given file based on the following
 - **Who:** A list of letters specifying whom you're going to be giving permissions to
 - ✓ **u** The user who owns the file (this means "you.")
 - ✓ **g** The group the file belongs to
 - ✓ **o** The other users
 - ✓ **a** all of the above (an abbreviation for **ugo**)
 - **Permissions**
 - ✓ **r** Permission to read the file
 - ✓ **w** Permission to write (or delete) the file
 - ✓ **x** Permission to execute the file, or, in the case of a directory, search it

For example, to:

- prevent outsiders from writing or executing file **archive.sh** one can issue command **chmod o=r archive.sh**
- take away all group permissions on **topsecret.inf** one can issue command **chmod g= topsecret.inf** (i.e. leave the permissions part of the command empty)
- allow reading and writing of **publicity.html** by anyone issue command **chmod og=rw publicity.html**
- (2) Change permissions on **wc.fields** so that the owner can **rw** while group and others can **r** (PS: you might need to run two different **chmod** commands)

E. Generalized Regular Expression Processor – grep: A great tool for searching through the contents of a file for fixed sequences of characters or regular expressions: e.g. **grep 'myString' Main.java**

- Searches for and displays all lines in file **Main.java** that contain the string **myString**
- The following flags may be used
 - **-i:** Ignore case differences
 - **-l:** only list filenames not actual lines
 - **-n:** show line numbers where matches were found
 - **-v:** reverse meaning of the search (i.e. all lines that don't match the pattern)
 - ✓ To search for all lines containing **println** except those using standard I/O (i.e. **System.out**)
 - ✓ **grep 'println' *.java | grep -v 'System.out'**

- (1) search for all lines (display line numbers as well) in file **story.txt** (created earlier) containing the string **'i'**
- To match a selection of characters, use **[]**: e.g. **[Hh]ello** matches lines containing **hello** or **Hello**
- Ranges of characters are also permitted

- [0-3] is the same as [0123]
- [a-k] is the same as [abcdefghijkl]
- [A-C] is the same as [ABC]
- [A-Ca-k] is the same as [ABCabcdefghijkl]

F. The find Command:

- E.g.: `find . -name '*JDBC*' -print` looks for any file whose name contains **JDBC**. It starts looking in the current directory (note the `.` in the command) and descends into all subdirectories. **-name** is called a predicate. It takes a regular expression as an argument. Any file that matches the predicate passes the control to the next predicate (**-print** in previous example)
- (1) find and display all java files (i.e. `-name '*.java'`) under your account [if it takes too long, press **CTRL-C** to cancel]
- The following command finds all files in folder **JDBC** (and descending from there) containing either string **JDBC** or **jDBC** in their names. Names of matching files are then passed to the **-exec** predicate which executes the next Linux command (i.e., `ls -l`) on them.
 - `find ./JDBC -name '*[jJ]DBC*.java' -exec ls -l '{}' \;`
 - PS: '{}' are replaced with the names of the matching files and `\;` indicates the end of the command (with no space between `\` and `;`)
- (2) change the above command and run it to display permissions on all files with names ending in `.java` under your account
- Other useful predicates
 - **-type d** → true if the file is a directory
 - **-type f** → true if the file is a plain file
 - **-mtime -5** → file is less than 5 days old (i.e. modified within the last 5 days ... **+5** means older than five days ... a 5 with no sign means exactly five days)
 - **-atime -5** → file was accessed less than 5 days ago
 - **-newer myEx.class** → file is newer than `myEx.class`
 - **-size +24k** → file is greater than 24k
- (3) Without running them, can you figure out what the following commands do?
 - `find . -name '*.java' -mtime -10 -atime +5 -print`
 - `find . -name '*.java' -mtime -10 -atime +5 -exec rm '{}' \;`